

## ▼ Overfitting and Underfitting

This notebook explores overfitting and underfitting scenarios based on [https://www.tensorflow.org/tutorials/keras/overfit\\_and\\_underfit](https://www.tensorflow.org/tutorials/keras/overfit_and_underfit)

## ▼ Setup and initialization

```
import tensorflow as tf

from tensorflow.keras import layers
from tensorflow.keras import regularizers

print(tf.__version__)

# let's set the random seed to make the results reproducible
tf.random.set_seed(74)
```

2.9.2

## ▼ Load dataset

```
#!pip install git+https://github.com/tensorflow/docs

import tensorflow_docs as tfdocs
import tensorflow_docs.modeling
import tensorflow_docs.plots

from IPython import display
from matplotlib import pyplot as plt

import numpy as np

import pathlib
import shutil
import tempfile
```

```
logdir = pathlib.Path(tempfile.mkdtemp())/ "tensorboard_logs"
shutil.rmtree(logdir, ignore_errors=True)
```

## ▼ The Higgs Dataset

```
gz = tf.keras.utils.get_file('HIGGS.csv.gz', 'http://mlphysics.ics.uci.edu/data/higgs/HIGGS.c
```

```
FEATURES = 28
```


```
ds = tf.data.experimental.CsvDataset(gz,[float(),]*(FEATURES+1), compression_type="GZIP")
```

```
def pack_row(*row):
    label = row[0]
    features = tf.stack(row[1:],1)
    return features, label
```

```
packed_ds = ds.batch(10000).map(pack_row).unbatch()
```

```
for features,label in packed_ds.batch(1000).take(1):
    print(features[0])
    plt.hist(features.numpy().flatten(), bins = 101)
```

```
tf.Tensor(
[ 0.8692932 -0.6350818  0.22569026  0.32747006 -0.6899932  0.75420225
 -0.24857314 -1.0920639  0.          1.3749921 -0.6536742  0.9303491
 1.1074361  1.1389043 -1.5781983 -1.0469854  0.          0.65792954
 -0.01045457 -0.04576717 3.1019614  1.35376  0.9795631  0.97807616
 0.92000484 0.72165745 0.98875093 0.87667835], shape=(28,), dtype=float32)
```



```
N_VALIDATION = int(1e3)
```

```

N_TRAIN = int(1e4)
BUFFER_SIZE = int(1e4)
BATCH_SIZE = 500
STEPS_PER_EPOCH = N_TRAIN//BATCH_SIZE

```

```

validate_ds = packed_ds.take(N_VALIDATION).cache()
train_ds = packed_ds.skip(N_VALIDATION).take(N_TRAIN).cache()

```

```
train_ds
```

```

<CacheDataset element_spec=(TensorSpec(shape=(28,), dtype=tf.float32, name=None),
TensorSpec(shape=(), dtype=tf.float32, name=None))>

```

```

validate_ds = validate_ds.batch(BATCH_SIZE)
train_ds = train_ds.shuffle(BUFFER_SIZE).repeat().batch(BATCH_SIZE)

```

## ▼ Demonstrating Overfitting

In deep learning, the number of learnable parameters in a model is often referred to as the model's "capacity".

Always keep this in mind: deep learning models tend to be good at fitting to the training data, but the real challenge is generalization, not fitting.

## ▼ Training procedure

Many models train better if you gradually reduce the learning rate during training. Use `tf.keras.optimizers.schedules` to reduce the learning rate over

```

lr_schedule = tf.keras.optimizers.schedules.InverseTimeDecay(
    0.001,
    decay_steps=STEPS_PER_EPOCH*1000,
    decay_rate=1,
    staircase=False)

```

```

def get_optimizer():
    return tf.keras.optimizers.Adam(lr_schedule)


```

```

step = np.linspace(0,100000)
lr = lr_schedule(step)
plt.figure(figsize = (8,6))
plt.plot(step/STEPS_PER_EPOCH, lr)
plt.ylim([0,max(plt.ylim())])

```

```
plt.xlabel('Epoch')
_ = plt.ylabel('Learning Rate')
```



Use `callbacks.TensorBoard` to generate TensorBoard logs for the training.

```
def get_callbacks(name):
    return [
        tfdocs.modeling.EpochDots(),
        tf.keras.callbacks.EarlyStopping(monitor='val_binary_crossentropy', patience=200),
        tf.keras.callbacks.TensorBoard(logdir/name),
    ]
```

Similarly each model will use the same `Model.compile` and `Model.fit` settings:

```
def compile_and_fit(model, name, optimizer=None, max_epochs=10000):
    if optimizer is None:
        optimizer = get_optimizer()
    model.compile(optimizer=optimizer,
                  loss=tf.keras.losses.BinaryCrossentropy(from_logits=True),
                  metrics=[
                      tf.keras.losses.BinaryCrossentropy(
                          from_logits=True, name='binary_crossentropy'),
                      'accuracy'])

    model.summary()

    history = model.fit(
        train_ds,
        steps_per_epoch = STEPS_PER_EPOCH,
```

```

epochs=max_epochs,
validation_data=validate_ds,
callbacks=get_callbacks(name),
verbose=0)
return history

```

## ▼ Tiny Model

```

tiny_model = tf.keras.Sequential([
    layers.Dense(16, activation='elu', input_shape=(FEATURES,)),
    layers.Dense(1)
])

```

```
size_histories = {}
```

```
size_histories['Tiny'] = compile_and_fit(tiny_model, 'sizes/Tiny')
```

```

.....^
Epoch: 2200, accuracy:0.6772, binary_crossentropy:0.5731, loss:0.5731, val_accuracy:0.6772
.....
Epoch: 2300, accuracy:0.6838, binary_crossentropy:0.5725, loss:0.5725, val_accuracy:0.6838
.....
Epoch: 2400, accuracy:0.6763, binary_crossentropy:0.5716, loss:0.5716, val_accuracy:0.6763
.....
Epoch: 2500, accuracy:0.6780, binary_crossentropy:0.5708, loss:0.5708, val_accuracy:0.6780
.....
Epoch: 2600, accuracy:0.6812, binary_crossentropy:0.5699, loss:0.5699, val_accuracy:0.6812
.....
Epoch: 2700, accuracy:0.6840, binary_crossentropy:0.5695, loss:0.5695, val_accuracy:0.6840
.....
Epoch: 2800, accuracy:0.6830, binary_crossentropy:0.5686, loss:0.5686, val_accuracy:0.6830
.....
Epoch: 2900, accuracy:0.6872, binary_crossentropy:0.5683, loss:0.5683, val_accuracy:0.6872
.....
Epoch: 3000, accuracy:0.6856, binary_crossentropy:0.5675, loss:0.5675, val_accuracy:0.6856
.....
Epoch: 3100, accuracy:0.6858, binary_crossentropy:0.5671, loss:0.5671, val_accuracy:0.6858
.....
Epoch: 3200, accuracy:0.6880, binary_crossentropy:0.5665, loss:0.5665, val_accuracy:0.6880
.....
Epoch: 3300, accuracy:0.6899, binary_crossentropy:0.5661, loss:0.5661, val_accuracy:0.6899
.....
Epoch: 3400, accuracy:0.6927, binary_crossentropy:0.5660, loss:0.5660, val_accuracy:0.6927
.....
Epoch: 3500, accuracy:0.6863, binary_crossentropy:0.5653, loss:0.5653, val_accuracy:0.6863
.....
Epoch: 3600, accuracy:0.6871, binary_crossentropy:0.5647, loss:0.5647, val_accuracy:0.6871
.....
Epoch: 3700, accuracy:0.6897, binary_crossentropy:0.5645, loss:0.5645, val_accuracy:0.6897
.....

```

```

.....
Epoch: 3800, accuracy:0.6919, binary_crossentropy:0.5641, loss:0.5641, val_accuracy:0.6898
.....
Epoch: 3900, accuracy:0.6894, binary_crossentropy:0.5637, loss:0.5637, val_accuracy:0.6898
.....
Epoch: 4000, accuracy:0.6900, binary_crossentropy:0.5633, loss:0.5633, val_accuracy:0.6898
.....
Epoch: 4100, accuracy:0.6887, binary_crossentropy:0.5631, loss:0.5631, val_accuracy:0.6898
.....
Epoch: 4200, accuracy:0.6929, binary_crossentropy:0.5626, loss:0.5626, val_accuracy:0.6898
.....
Epoch: 4300, accuracy:0.6930, binary_crossentropy:0.5621, loss:0.5621, val_accuracy:0.6898
.....
Epoch: 4400, accuracy:0.6920, binary_crossentropy:0.5619, loss:0.5619, val_accuracy:0.6898
.....
Epoch: 4500, accuracy:0.6913, binary_crossentropy:0.5617, loss:0.5617, val_accuracy:0.6898
.....
Epoch: 4600, accuracy:0.6945, binary_crossentropy:0.5619, loss:0.5619, val_accuracy:0.6898
.....
Epoch: 4700, accuracy:0.6934, binary_crossentropy:0.5611, loss:0.5611, val_accuracy:0.6898
.....
Epoch: 4800, accuracy:0.6966, binary_crossentropy:0.5609, loss:0.5609, val_accuracy:0.6898
.....
Epoch: 4900, accuracy:0.6963, binary_crossentropy:0.5607, loss:0.5607, val_accuracy:0.6898
.....
Epoch: 5000, accuracy:0.6898, binary_crossentropy:0.5605, loss:0.5605, val_accuracy:0.6898


```

Now check how the model did:

```

plotter = tfdocs.plots.HistoryPlotter(metric = 'binary_crossentropy', smoothing_std=10)
plotter.plot(size_histories)
plt.ylim([0.5, 0.7])

```



▼ Small model

To check if you can beat the performance of the small model, progressively train some larger models.

Try two hidden layers with 16 units each:

```
small_model = tf.keras.Sequential([
    # `input_shape` is only required here so that `.summary` works.
    layers.Dense(32, activation='elu', input_shape=(FEATURES,)),
    layers.Dense(32, activation='elu'),
    layers.Dense(1)
])
```

```
size_histories['Small'] = compile_and_fit(small_model, 'sizes/Small')
```

Model: "sequential\_6"


Layer (type)	Output Shape	Param #
dense_21 (Dense)	(None, 32)	928
dense_22 (Dense)	(None, 32)	1056
dense_23 (Dense)	(None, 1)	33

=====  
Total params: 2,017  
Trainable params: 2,017  
Non-trainable params: 0

Epoch: 0, accuracy:0.4793, binary\_crossentropy:0.7203, loss:0.7203, val\_accuracy:0.4  
.....  
Epoch: 100, accuracy:0.6156, binary\_crossentropy:0.6113, loss:0.6113, val\_accuracy:0  
.....  
Epoch: 200, accuracy:0.6791, binary\_crossentropy:0.5718, loss:0.5718, val\_accuracy:0  
.....  
Epoch: 300, accuracy:0.6984, binary\_crossentropy:0.5479, loss:0.5479, val\_accuracy:0  
.....  
Epoch: 400, accuracy:0.7145, binary\_crossentropy:0.5349, loss:0.5349, val\_accuracy:0  
.....



```
# plot
#plotter = tfdocs.plots.HistoryPlotter(metric = 'binary_crossentropy', smoothing_std=10)
plotter.plot(size_histories)
plt.ylim([0.5, 0.7])
```



## ▼ Medium model

```
medium_model = tf.keras.Sequential([
    layers.Dense(64, activation='elu', input_shape=(FEATURES,)),
    layers.Dense(64, activation='elu'),
    layers.Dense(64, activation='elu'),
    layers.Dense(1)
])
```

And train the model using the same data:

```
size_histories['Medium'] = compile_and_fit(medium_model, "sizes/Medium")
```

Model: "sequential\_7"

Layer (type)	Output Shape	Param #
dense_24 (Dense)	(None, 64)	1856
dense_25 (Dense)	(None, 64)	4160
dense_26 (Dense)	(None, 64)	4160
dense_27 (Dense)	(None, 1)	65

```
=====
Total params: 10,241
Trainable params: 10,241
Non-trainable params: 0
=====
```

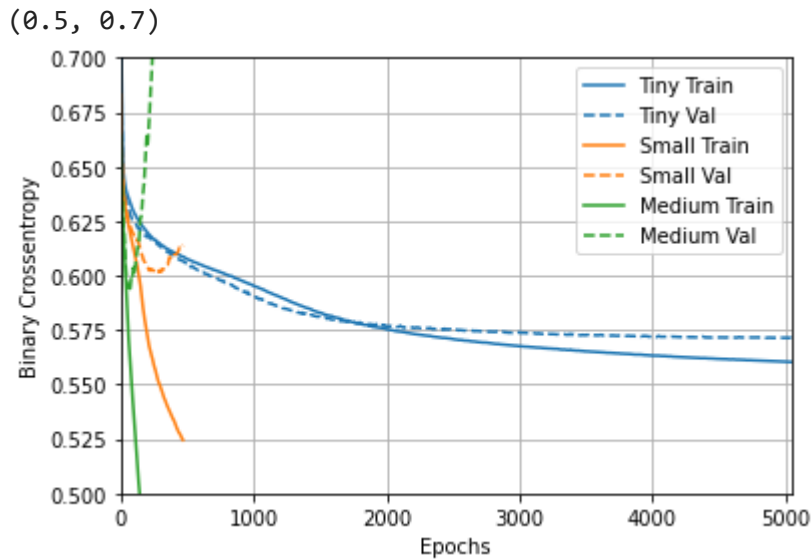
```
Epoch: 0, accuracy:0.4973, binary_crossentropy:0.7047, loss:0.7047, val_accuracy:0.4
.....
Epoch: 100, accuracy:0.7171, binary_crossentropy:0.5324, loss:0.5324, val_accuracy:0
.....
```



Epoch: 200, accuracy:0.7767, binary\_crossentropy:0.4458, loss:0.4458, val\_accuracy:0  
.....



```
plotter.plot(size_histories)
plt.ylim([0.5, 0.7])
```



## ▼ Large model

```
large_model = tf.keras.Sequential([
    layers.Dense(512, activation='elu', input_shape=(FEATURES,)),
    layers.Dense(512, activation='elu'),
    layers.Dense(512, activation='elu'),
    layers.Dense(512, activation='elu'),
    layers.Dense(1)
])
```

And, again, train the model using the same data:

```
size_histories['large'] = compile_and_fit(large_model, "sizes/large")
```

Model: "sequential\_8"

Layer (type)	Output Shape	Param #
dense_28 (Dense)	(None, 512)	14848
dense_29 (Dense)	(None, 512)	262656
dense_30 (Dense)	(None, 512)	262656


dense_31 (Dense)	(None, 512)	262656
dense_32 (Dense)	(None, 1)	513

```
=====
Total params: 803,329
Trainable params: 803,329
Non-trainable params: 0
=====
```

```
Epoch: 0, accuracy:0.5066, binary_crossentropy:0.8190, loss:0.8190, val_accuracy:0.4
.....
Epoch: 100, accuracy:1.0000, binary_crossentropy:0.0020, loss:0.0020, val_accuracy:0
.....
Epoch: 200, accuracy:1.0000, binary_crossentropy:0.0001, loss:0.0001, val_accuracy:0
.....
```



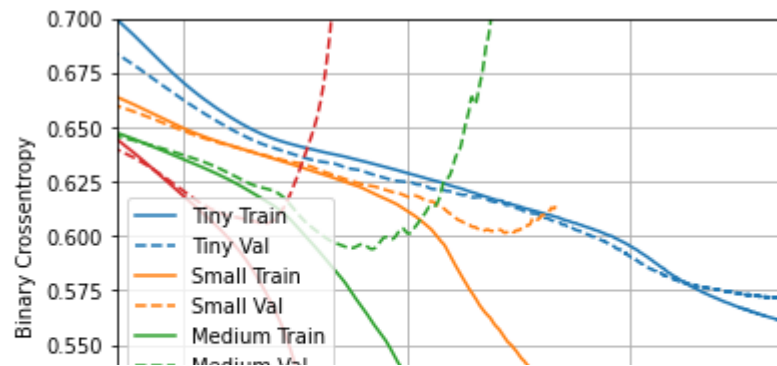
```
plotter.plot(size_histories)
plt.ylim([0.5, 0.7])
```



## ▼ Plot the training and validation losses

```
plotter.plot(size_histories)
a = plt.xscale('log')
plt.xlim([5, max(plt.xlim())])
plt.ylim([0.5, 0.7])
plt.xlabel("Epochs [Log Scale]")
```

Text(0.5, 0, 'Epochs [Log Scale]')



## ▼ View in TensorBoard

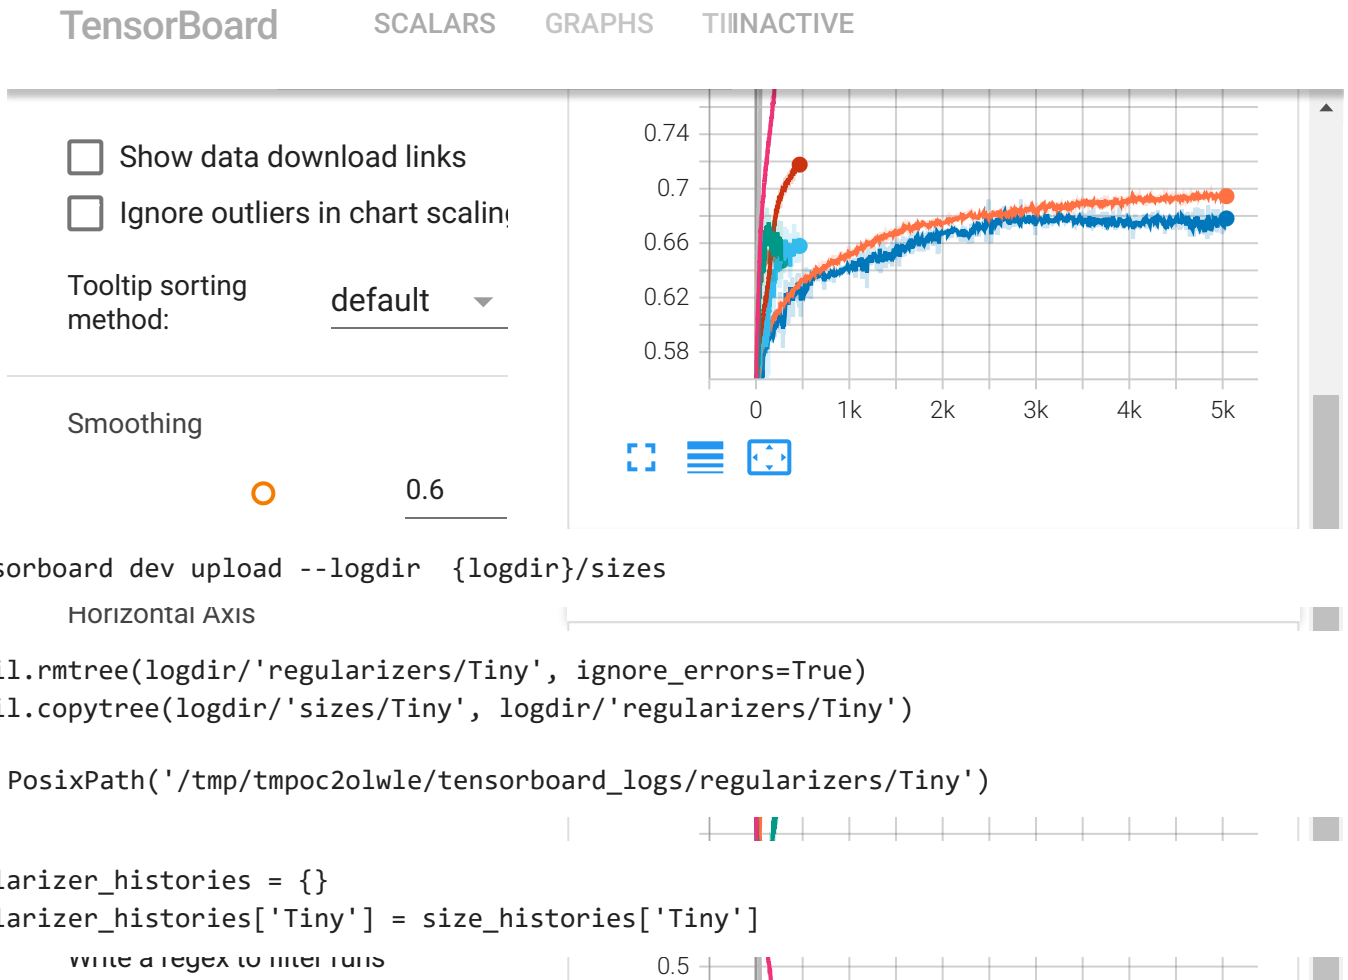
Epochs [Log Scale]

```
#docs_infra: no_execute
```

```
# Load the TensorBoard notebook extension  
%load_ext tensorboard
```

```
# Open an embedded TensorBoard viewer  
%tensorboard --logdir {logdir}/sizes
```

The tensorboard extension is already loaded. To reload it, use:  
`%reload_ext tensorboard`



## ▼ Add weight regularization

L1 regularization, where the cost added is proportional to the absolute value of the weights coefficients (i.e. to what is called the "L1 norm" of the weights).

L2 regularization, where the cost added is proportional to the square of the value of the weights coefficients (i.e. to what is called the squared "L2 norm" of the weights). L2 regularization is also called weight decay in the context of neural networks. Don't let the different name confuse you: weight decay is mathematically the exact same as L2 regularization.

L1 regularization pushes weights towards exactly zero, encouraging a sparse model. L2 regularization will penalize the weights parameters without making them sparse since the penalty goes to zero for small weights—one reason why L2 is more common.

In `tf.keras`, weight regularization is added by passing weight regularizer instances to layers as keyword arguments. Add L2 weight regularization:

```
l2_model = tf.keras.Sequential([
```

```

layers.Dense(512, activation='elu',
              kernel_regularizer=regularizers.l2(0.001),
              input_shape=(FEATURES,)),
layers.Dense(512, activation='elu',
              kernel_regularizer=regularizers.l2(0.001)),
layers.Dense(512, activation='elu',
              kernel_regularizer=regularizers.l2(0.001)),
layers.Dense(512, activation='elu',
              kernel_regularizer=regularizers.l2(0.001)),
layers.Dense(1)
])

```

```

regularizer_histories['l2'] = compile_and_fit(l2_model, "regularizers/l2")

```

Model: "sequential\_9"

Layer (type)	Output Shape	Param #
dense_33 (Dense)	(None, 512)	14848
dense_34 (Dense)	(None, 512)	262656
dense_35 (Dense)	(None, 512)	262656
dense_36 (Dense)	(None, 512)	262656
dense_37 (Dense)	(None, 1)	513

```

=====
Total params: 803,329
Trainable params: 803,329
Non-trainable params: 0
=====

```

```

Epoch: 0, accuracy:0.5080, binary_crossentropy:0.7454, loss:2.2338, val_accuracy:0.4
.....
Epoch: 100, accuracy:0.6524, binary_crossentropy:0.5976, loss:0.6200, val_accuracy:0
.....
Epoch: 200, accuracy:0.6715, binary_crossentropy:0.5849, loss:0.6074, val_accuracy:0
.....
Epoch: 300, accuracy:0.6794, binary_crossentropy:0.5823, loss:0.6087, val_accuracy:0
.....
Epoch: 400, accuracy:0.6945, binary_crossentropy:0.5616, loss:0.5878, val_accuracy:0
.....
Epoch: 500, accuracy:0.6944, binary_crossentropy:0.5574, loss:0.5842, val_accuracy:0
.....
Epoch: 600, accuracy:0.7053, binary_crossentropy:0.5484, loss:0.5758, val_accuracy:0
.....
Epoch: 700, accuracy:0.7098, binary_crossentropy:0.5411, loss:0.5687, val_accuracy:0
.....

```




```

plotter.plot(regularizer_histories)
plt.ylim([0.5, 0.7])

```

(0.5, 0.7)



[Colab paid products](#) - [Cancel contracts here](#)

✓ 12s completed at 11:53 AM

